

Assignment 2

Deadline: Sunday, February 20, at 11:59pm.

Submission: You need to submit the final PDF file, and any Python scripts, as a compressed folder (.zip or .tgz) on Quercus. If you used Google Colab, including a link is sufficient.

Neatness Point: One point will be given for neatness. You will receive this point as long as we don't have a hard time reading your solutions or understanding the structure of your code.

Late Submission: 10% of the marks will be deducted for each day late, up to a maximum of 3 days. After that, no submissions will be accepted.

Assignments are individual work. See the Course Information handout¹ for detailed policies.

1. [2pts] Decision Trees for Heart Failure Prediction.

The following table contains training examples that help predict whether a patient is likely to develop a heart failure.

PATIENT ID	CHEST PAIN	MALE	SMOKES	EXERCISES	HEART FAILURE
1	1	1	0	1	1
2	1	1	1	0	1
3	0	0	1	0	1
4	0	1	0	1	0
5	1	0	1	1	1
6	0	1	1	1	0

Table 1: Training Data for Decision Trees

Use information gain to construct a minimal decision tree that predicts whether or not a patient is likely to have a heart failure. **SHOW EACH STEP OF THE COMPUTATION.**

2. [3pts] Revisiting cell classification with Random Forests.

In Assignment 1 Q4, you applied KNN to classify benign and malignant cells from image features. In this question, you will train a random forest model on the same dataset, and compare its results to the simpler KNN approach.

- (a) [1pts] Use `sklearn.ensemble.RandomForestClassifier` to train a random forest classifier. Similar to how you varied the value of `k` in 4b from A1, you should try different values of `n_estimators` (10, 20, 40, 100) and make plots of training and validation accuracy. Then, report the test accuracy for the best model. Is your best Random Forest better than your best KNN classifier?
- (b) [2pts] One particularly useful attribute of Random Forests is that they provide a simple way to interpret feature importance. By default, the fitted `RandomForestClassifier`

¹<https://lmp1210-uoft.github.io/2022/assets/misc/syllabus.pdf>

provides a method, `feature_importances_`, which gives the impurity-based feature importance. Scikit-learn provides an alternate approach, permutation-based feature importance, through the function `sklearn.inspection.permutation_importance`. Plot the top 5 and bottom 5 features for your best-performing model using each method. For `permutation_importance`, use `scoring="accuracy"` and the training set as the data (to make it consistent with the other method), and make sure to use the `importances_mean` of the returned dictionary. Do the methods agree on which features are important?

3. [3pts] Binary Classification with Imbalanced Data

In Lecture 2 we talked about different binary classification metrics and their use cases. In this problem we will illustrate some of the strengths and limitations of these metrics.

- (a) [1pts] Scikit-learn provides tools for generating synthetic ("toy") datasets. Use the `sklearn.datasets.make_classification` function to construct a dataset, with `n_samples=1000`, `n_features=20`, and `weights=[0.90,0.1]`. This will produce a dataset where the inputs are 10-dimensional vectors and labels are binary. Note that the dataset is highly imbalanced, with 90% negative examples (0s) and 10% positive examples (1s). Make a train/test split using `sklearn.model_selection.train_test_split`, where 70% of the data is for training and 30% of the data is for testing².

Train a logistic regression using `sklearn.neural_network.LogisticRegression` (with default parameters) on this data, and report the model's accuracy, precision, recall, and F1 score. In addition, plot the ROC and PR curves, using `sklearn.metrics.precision_recall_curve` and `sklearn.metrics.roc_curve`.

Bonus (no credit): try implementing the scoring functions yourself!

- (b) [2pts] Propose one way to improve your model's recall without ruining its precision. Feel free to look up common solutions to class imbalance online. Implement this change and train a new logistic regression model. Report the same metrics as in part (a) (including the plots). Were you able to improve the recall, and how much precision did it cost you?

4. [2pts] Iterative and Analytic Solutions of Linear Regression

Linear regression has an analytic solution, but it can also be solved optimally using gradient descent (with the appropriate learning rate). Using the starter code in this [Google Colab](#), implement the analytic solution for linear regression and the solution that uses (full batch) gradient descent. For simplicity, this model does not have any bias term. You can submit a link to your completed Colab book as the answer to this question.

5. [2pts] **Gradient Descent** We have a model which takes a two-dimensional input $x = (x_1, x_2)$ and has two parameters $w = (w_1, w_2)$ given by $f(x, w) = \sigma(\sigma(x_1 w_1) w_2 + x_2)$ where $\sigma(z) = \frac{1}{1+e^{-z}}$. Now we use backpropagation to estimate the right parameter values. We start by setting both the parameters to 0. Assume that we are given a training point $\mathbf{x}_1 = (1, 0)$ and its target (or label) $y = 5$, what is the value of the gradient $\frac{\partial f}{\partial w_2}$? If the learning rate is 0.5, what will be the value of w_2 after one update using gradient descent algorithm if we use the loss function $L = (f - y)^2$? **PLEASE SHOW ALL THE STEPS!**

²for our purposes, we are not playing with any of the model hyperparameters, so we don't need a separate validation set for hyperparameter selection

6. [2pts] **Computation Graph and Backpropagation** Assume we have a simple function $f(w, x, y, z) = (x - y)^2 * (w + \exp(z))$. Please draw the corresponding computation graph and derive the following four derivatives: $\frac{\partial f}{\partial w}$, $\frac{\partial f}{\partial x}$, $\frac{\partial f}{\partial y}$, $\frac{\partial f}{\partial z}$ using backpropagation.